


Shown: 1 of 10 identical members — each differs only by init seed. Conv(k×k, pad k/2, no bias) + BatchNorm + ReLU per block; height ∝ channel width; grid overlay = 3×3 kernel window.

10 members · seeds 42–51
981K params each · kernel 3×3



■ Conv 3×3 + BN + ReLU (stack depth = channel count)

□ feature map / kernel window overlay

$\times \frac{1}{2}$ = MaxPool 2×2 (all blocks but the last)

----- frozen-trunk decoder tap (ConvTranspose 4×4 s2 stages, reconstructs 128^2 input — drives ensemble bias/variance decomposition)

ENSEMBLE MEMBERS (IDENTICAL ARCHITECTURE, CHANNELS 32-64-128-256-256)

- Model 1 · seed 42

● Model 2 · seed 43

● Model 3 · seed 44

● Model 4 · seed 45

● Model 5 · seed 46
- Model 6 · seed 47

● Model 7 · seed 48

● Model 8 · seed 49

● Model 9 · seed 50

● Model 10 · seed 51