


Each block: Conv(k×k, pad k/2, no bias) + BatchNorm + ReLU; $\times\frac{1}{2}$ = MaxPool 2×2 (all blocks but the last). Rectangle height ∝ channel width; colour = kernel size.



kernel 3×3 kernel 5×5 decoder tap

dec Lk: frozen-trunk BlockDecoder (ConvTranspose 4×4 stride-2 stages + BN + ReLU) reconstructing the 128² input from block k — the reconstructions drive the ensemble bias/variance decomposition.