

Each cube: Conv(k×k, pad k/2, no bias) + BatchNorm + ReLU, then 2×2 max-pool (all stages but the last). Cube height ∝ channel width; colour = kernel size — the shape of each staircase is where the members diverge.



kernel 3×3 kernel 5×5 cube height ∝ channel count (16 – 384)

All members share the same prediction heads on a global-average-pooled trunk (gender → 2 · attrs → 6 · cutpaste → 2). Each stage also feeds a frozen-trunk decoder (ConvTranspose 4×4 stride-2 + BN + ReLU) reconstructing the 128³ input from that layer — used for the ensemble's bias/variance decomposition.