


Each row = one decoder block. ConvTranspose2d(4×4, s2, p1) → BatchNorm → ReLU repeated per upsample step; channels halve each step, floor 16. A 1×1 head + Sigmoid produces the block's own RGB prediction.



— sequential step chip chain = upsample stages within the block (channels halve, floor 16)

1,615,391 params · lrad/decoder.py · lrad/train.py:train_decoders